

Design Considerations for PL/SQL Code

ORACLE

Copyright © 2019, Oracle and/or its affiliates. All rights reserved.

Objectives

After completing this lesson, you should be able to do the following:

- Using the RETURNING Clause
- Using Bulk Binding
- Bulk Binding: FORALL Keyword
- Bulk Binding: BULK COLLECT Keyword
- Bulk Binding: Syntax and Keywords
- Using BULK COLLECT INTO with Queries
- Using BULK COLLECT INTO with Cursors

Using the RETURNING Clause

- Often, applications need information about the row affected by a SQL operation.
- e.g; To generate a report or to take a subsequent action.
- The `INSERT`, `UPDATE`, and `DELETE` statements can include a `RETURNING` clause, which returns column values from the affected row into PL/SQL variables or host variables.
- This eliminates the need to `SELECT` the row after an `INSERT` or `UPDATE`, or before a `DELETE`.
- As a result, fewer network round trips, less server CPU time, fewer cursors, and less server memory are required.

Using the RETURNING Clause

- Improves performance by returning column values with INSERT, UPDATE, and DELETE statements
- Eliminates the need for a SELECT statement

```
CREATE OR REPLACE PROCEDURE update_salary(p_emp_id
    NUMBER) IS
    v_name      employees.last_name%TYPE;
    v_new_sal    employees.salary%TYPE;
BEGIN
    UPDATE employees
        SET salary = salary * 1.1
    WHERE employee_id = p_emp_id
    RETURNING last_name, salary INTO v_name, v_new_sal;
    DBMS_OUTPUT.PUT_LINE(v_name || ' new salary is ' ||
        v_new_sal);
END update_salary;
/
```

ORACLE

```
1 SET SERVEROUTPUT ON
2 /
3 SELECT last_name, salary
4 FROM employees
5 WHERE employee_id = 108;
6 /
7 EXECUTE update_salary(108)
```

LAST_NAME	SALARY
Greenberg	12008

anonymous block completed
Greenberg new salary is 13208.8

Using Bulk Binding

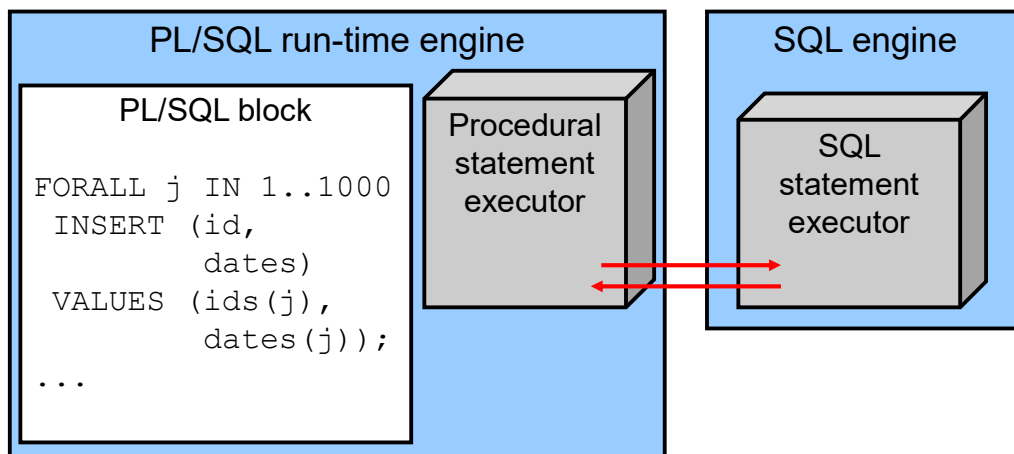
- The Oracle server uses two engines to run PL/SQL blocks and subprograms:
 - The PL/SQL run-time engine, which runs procedural statements but passes the SQL statements to the SQL engine
 - The SQL engine, which parses and executes the SQL statement and, in some cases, returns data to the PL/SQL engine
- During execution, every SQL statement causes a context switch between the two engines, which results in a performance penalty for excessive amounts of SQL processing.
- This is typical of applications that have a SQL statement in a loop that uses values from indexed collections.
- Collections include nested tables, varying arrays, index-by tables, and host arrays.

Using Bulk Binding

- Performance can be substantially improved by minimizing the number of context switches through the use of **bulk binding**.
- Bulk binding causes an entire collection to be bound in one call, a context switch, to the SQL engine.
- That is, a bulk bind process passes the entire collection of values back and forth between the two engines in a single context switch, compared with incurring a context switch for each collection element in an iteration of a loop.
- The more rows affected by a SQL statement, the greater the performance gain with bulk binding.

Using Bulk Binding

Binds whole arrays of values in a single operation, rather than using a loop to perform a `FETCH`, `INSERT`, `UPDATE`, and `DELETE` operation multiple times



Bulk Binding: FORALL Keyword

- Use bulk binds to improve the performance of:
 - DML statements that reference collection elements
 - `SELECT` statements that reference collection elements
 - Cursor `FOR` loops that reference collections and the `RETURNING INTO` clause
- The `FORALL` keyword instructs the PL/SQL engine to bulk bind input collections before sending them to the SQL engine.
- Although the `FORALL` statement contains an iteration scheme, it is not a `FOR` loop.

Bulk Binding: BULK COLLECT Keyword

- The `BULK COLLECT` keyword instructs the SQL engine to bulk bind output collections, before returning them to the PL/SQL engine.
- This enables you to bind locations into which SQL can return the retrieved values in bulk.
- Thus, you can use these keywords in the `SELECT INTO`, `FETCH INTO`, and `RETURNING INTO` clauses.

Bulk Binding: Syntax and Keywords

- The `FORALL` keyword instructs the *PL/SQL engine* to bulk bind input collections before sending them to the SQL engine.

```
FORALL index IN lower_bound .. upper_bound  
  [SAVE EXCEPTIONS]  
  sql_statement;
```

- The `BULK COLLECT` keyword instructs the *SQL engine* to bulk bind output collections before returning them to the PL/SQL engine.

```
... BULK COLLECT INTO  
    collection_name[,collection_name] ...
```

ORACLE

The `SAVE EXCEPTIONS` keyword is optional. However, if some of the DML operations succeed and some fail, you would want to track or report on those that fail. Using the `SAVE EXCEPTIONS` keyword causes failed operations to be stored in a cursor attribute called `%BULK_EXCEPTIONS`, which is a collection of records indicating the bulk DML iteration number and corresponding error code.

Handling FORALL Exceptions with the %BULK_EXCEPTIONS Attribute

To manage exceptions and have the bulk bind complete despite errors, add the keywords `SAVE EXCEPTIONS` to your `FORALL` statement after the bounds, before the DML statement.

All exceptions raised during the execution are saved in the cursor attribute `%BULK_EXCEPTIONS`, which stores a collection of records. Each record has two fields:

`%BULK_EXCEPTIONS(i).ERROR_INDEX` holds the “iteration” of the `FORALL` statement during which the exception was raised and `%BULK_EXCEPTIONS(i).ERROR_CODE` holds the corresponding Oracle error code.

Values stored in `%BULK_EXCEPTIONS` refer to the most recently executed `FORALL` statement. Its subscripts range from 1 to `%BULK_EXCEPTIONS.COUNT`.

Note: For additional information about bulk binding and handling bulk-binding exceptions, refer to the *Oracle Database PL/SQL User's Guide and Reference*.

Bulk Binding FORALL: Example

```
CREATE PROCEDURE raise_salary(p_percent NUMBER) IS
  TYPE numlist_type IS TABLE OF NUMBER
    INDEX BY BINARY_INTEGER;
  v_id numlist_type; -- collection
BEGIN
  v_id(1) := 100; v_id(2) := 102; v_id(3) := 104; v_id(4) := 110;
  -- bulk-bind the PL/SQL table
  FORALL i IN v_id.FIRST .. v_id.LAST
    UPDATE employees
      SET salary = (1 + p_percent/100) * salary
      WHERE employee_id = v_id(i);
END;
/
```

```
EXECUTE raise_salary(10)
```

anonymous block completed

ORACLE

Note: Before you can run the example in the slide, you must disable the `update_job_history` trigger as follows:

```
ALTER TRIGGER update_job_history DISABLE;
```

In the example in the slide, the PL/SQL block increases the salary for employees with IDs 100, 102, 104, or 110. It uses the `FORALL` keyword to bulk bind the collection. Without bulk binding, the PL/SQL block would have sent a SQL statement to the SQL engine for each employee record that is updated. If there are many employee records to update, the large number of context switches between the PL/SQL engine and the SQL engine can affect performance. However, the `FORALL` keyword bulk binds the collection to improve performance.

Note: A looping construct is no longer required when using this feature.

An Additional Cursor Attribute for DML Operations

Another cursor attribute added to support bulk operations is %BULK_ROWCOUNT. The %BULK_ROWCOUNT attribute is a composite structure designed for use with the FORALL statement. This attribute acts like an index-by table. Its *i*th element stores the number of rows processed by the *i*th execution of an UPDATE or DELETE statement. If the *i*th execution affects no rows, then %BULK_ROWCOUNT(*i*) returns zero.

Here is an example:

```
CREATE TABLE num_table (n NUMBER);
DECLARE
    TYPE num_list_type IS TABLE OF NUMBER
        INDEX BY BINARY_INTEGER;
    v_nums num_list_type;
BEGIN
    v_nums(1) := 1;
    v_nums(2) := 3;
    v_nums(3) := 5;
    v_nums(4) := 7;
    v_nums(5) := 11;
    FORALL i IN v_nums.FIRST .. v_nums.LAST
        INSERT INTO v_num_table (n) VALUES (v_nums(i));
    FOR i IN v_nums.FIRST .. v_nums.LAST
    LOOP
        dbms_output.put_line('Inserted ' ||
            SQL%BULK_ROWCOUNT(i) || ' row(s) '
            || ' on iteration ' || i);
    END LOOP;
END;
/
DROP TABLE num_table;
```

The following results are produced by this example:

```
table NUM_TABLE created.
anonymous block completed
Inserted 1 row(s) on iteration 1
Inserted 1 row(s) on iteration 2
Inserted 1 row(s) on iteration 3
Inserted 1 row(s) on iteration 4
Inserted 1 row(s) on iteration 5

table NUM_TABLE dropped.
```

Using BULK COLLECT INTO with Queries

The **SELECT** statement supports the **BULK COLLECT INTO** syntax.

```
CREATE PROCEDURE get_departments(p_loc NUMBER) IS
  TYPE dept_tab_type IS
    TABLE OF departments%ROWTYPE;
  v_depts dept_tab_type;
BEGIN
  SELECT * BULK COLLECT INTO v_depts
  FROM departments
  WHERE location_id = p_loc;
  FOR i IN 1 .. v_depts.COUNT LOOP
    DBMS_OUTPUT.PUT_LINE(v_depts(i).department_id
      || ' ' || v_depts(i).department_name);
  END LOOP;
END;
```

ORACLE

Starting with Oracle Database 10g, when using a **SELECT** statement in PL/SQL, you can use the bulk collection syntax shown in the example in the slide. Thus, you can quickly obtain a set of rows without using a cursor mechanism.

The example reads all the department rows for a specified region into a PL/SQL table, whose contents are displayed with the **FOR** loop that follows the **SELECT** statement.

Using BULK COLLECT INTO with Cursors

The **FETCH** statement has been enhanced to support the **BULK COLLECT INTO** syntax.

```
CREATE OR REPLACE PROCEDURE get_departments(p_loc NUMBER) IS
  CURSOR cur_dept IS
    SELECT * FROM departments
    WHERE location_id = p_loc;
  TYPE dept_tab_type IS TABLE OF cur_dept%ROWTYPE;
  v_depts dept_tab_type;
BEGIN
  OPEN cur_dept;
  FETCH cur_dept BULK COLLECT INTO v_depts;
  CLOSE cur_dept;
  FOR i IN 1 .. v_depts.COUNT LOOP
    DBMS_OUTPUT.PUT_LINE(v_depts(i).department_id
      || ' ' || v_depts(i).department_name);
  END LOOP;
END;
```

Quiz

The `NOCOPY` hint allows the PL/SQL compiler to pass `OUT` and `IN OUT` parameters by reference rather than by value. This enhances performance by reducing overhead when passing parameters.

- a. True
- b. False

ORACLE

Answer: a

PL/SQL subprograms support three parameter-passing modes: `IN`, `OUT`, and `IN OUT`.

By default:

- The `IN` parameter is passed by reference. A pointer to the `IN` actual parameter is passed to the corresponding formal parameter. So, both the parameters reference the same memory location, which holds the value of the actual parameter.
- The `OUT` and `IN OUT` parameters are passed by value. The value of the `OUT` or `IN OUT` actual parameter is copied into the corresponding formal parameter. Then, if the subprogram exits normally, the values assigned to the `OUT` and `IN OUT` formal parameters are copied into the corresponding actual parameters.

Copying parameters that represent large data structures (such as collections, records, and instances of object types) with `OUT` and `IN OUT` parameters slows down execution and uses up memory. To prevent this overhead, you can specify the `NOCOPY` hint, which enables the PL/SQL compiler to pass `OUT` and `IN OUT` parameters by reference.

Summary

In this lesson, you should have learned how to:

- Using the RETURNING Clause
- Using Bulk Binding
- Bulk Binding: FORALL Keyword
- Bulk Binding: BULK COLLECT Keyword
- Bulk Binding: Syntax and Keywords
- Using BULK COLLECT INTO with Queries
- Using BULK COLLECT INTO with Cursors